

HOTEL RESERVATION CANCELLATION PREDICTION

Data Description

The online hotel reservation channels have dramatically changed booking possibilities and customers' behavior. A significant number of hotel reservations are called-off due to cancellations or no-shows. The typical reasons for cancellations include change of plans, scheduling conflicts, etc. This is often made easier by the option to do so free of charge or preferably at a low cost which is beneficial to hotel guests but it is a less desirable and possibly revenue-diminishing factor for hotels to deal with.

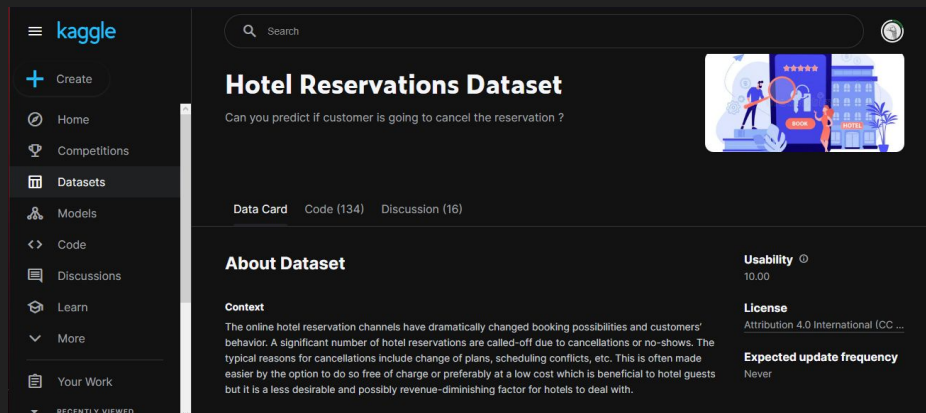
The Goal

Predicting if the customer is going to honor the reservation or cancel it

SOURCE OF DATA

The dataset used for this project is taken from Kaggle with the name 'Hotel Reservations Dataset' which can be accessed through this link:

<https://www.kaggle.com/datasets/ahsan81/hotel-reservations-classification-dataset>



The screenshot shows the Kaggle interface for the 'Hotel Reservations Dataset'. On the left is a navigation sidebar with links to Home, Competitions, Datasets (selected), Models, Code, Discussions, Learn, More, and Your Work. The main content area features a search bar, the dataset title 'Hotel Reservations Dataset', and a description: 'Can you predict if customer is going to cancel the reservation?'. Below this are tabs for 'Data Card' (selected), 'Code (134)', and 'Discussion (16)'. The 'About Dataset' section includes a 'Context' paragraph about online hotel reservation channels and booking possibilities. On the right, there are three metrics: 'Usability' (10.00), 'License' (Attribution 4.0 International (CC BY-SA 4.0)), and 'Expected update frequency' (Never). An illustration of a person at a computer with a hotel building in the background is also visible.

kaggle

+ Create

Home

Competitions

Datasets

Models

Code

Discussions

Learn

More

Your Work

RECENTLY VIEWED

Search

Hotel Reservations Dataset

Can you predict if customer is going to cancel the reservation ?

Data Card Code (134) Discussion (16)

About Dataset

Context

The online hotel reservation channels have dramatically changed booking possibilities and customers' behavior. A significant number of hotel reservations are called-off due to cancellations or no-shows. The typical reasons for cancellations include change of plans, scheduling conflicts, etc. This is often made easier by the option to do so free of charge or preferably at a low cost which is beneficial to hotel guests but it is a less desirable and possibly revenue-diminishing factor for hotels to deal with.

Usability 10.00

License Attribution 4.0 International (CC BY-SA 4.0)

Expected update frequency Never

IMPORTING PACKAGES AND DATASET

```
In [1]: import numpy as np  
import pandas as pd
```

```
In [2]: import matplotlib.pyplot as plt  
import seaborn as sns
```



PACKAGES USED

```
data = pd.read_csv('Hotel Reservations.csv')  
data.head()
```

	Booking_ID	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_nights	type_of_meal_plan	req
0	INN00001	2	0	1	2	Meal Plan 1	0
1	INN00002	2	0	2	3	Not Selected	0
2	INN00003	1	0	2	1	Meal Plan 1	0
3	INN00004	2	0	0	2	Meal Plan 1	0
4	INN00005	2	0	1	1	Not Selected	0



IMPORTING
DATASET

DATA PROFILING

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 36275 entries, 0 to 36274  
Data columns (total 19 columns):
```

#	Column	Non-Null Count	Dtype
0	Booking_ID	36275 non-null	object
1	no_of_adults	36275 non-null	int64
2	no_of_children	36275 non-null	int64
3	no_of_weekend_nights	36275 non-null	int64
4	no_of_week_nights	36275 non-null	int64
5	type_of_meal_plan	36275 non-null	object
6	required_car_parking_space	36275 non-null	int64
7	room_type_reserved	36275 non-null	object
8	lead_time	36275 non-null	int64
9	arrival_year	36275 non-null	int64
10	arrival_month	36275 non-null	int64
11	arrival_date	36275 non-null	int64
12	market_segment_type	36275 non-null	object
13	repeated_guest	36275 non-null	int64
14	no_of_previous_cancellations	36275 non-null	int64
15	no_of_previous_bookings_not_canceled	36275 non-null	int64
16	avg_price_per_room	36275 non-null	float64
17	no_of_special_requests	36275 non-null	int64
18	booking_status	36275 non-null	object

```
dtypes: float64(1), int64(13), object(5)  
memory usage: 5.3+ MB
```

This dataset consists of 36275 rows and 19 columns.

```
# Checking for missing values
```


```
data.isnull().sum()
```

Booking_ID	0
no_of_adults	0
no_of_children	0
no_of_weekend_nights	0
no_of_week_nights	0
type_of_meal_plan	0
required_car_parking_space	0
room_type_reserved	0
lead_time	0
arrival_year	0
arrival_month	0
arrival_date	0
market_segment_type	0
repeated_guest	0
no_of_previous_cancellations	0
no_of_previous_bookings_not_canceled	0
avg_price_per_room	0
no_of_special_requests	0
booking_status	0
dtype: int64	

There are no missing values in this dataset.

PROFILING DESCRIPTIVE STATISTICS

```
# Distribution of each features  
data.hist()  
plt.show()
```



PROFILING DESCRIPTIVE STATISTICS

	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_nights	required_car_parking_space	lead_time	arrival_year	arrival_month	arrival_date	repeated_guest	no_of_previous_cancellations
count	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000	36275.000000
mean	1.844962	0.105279	0.810724	2.204300	0.030986	85.232557	2017.820427	7.423653	15.596995	0.025637	0.023349
std	0.518715	0.402648	0.870644	1.410905	0.173281	85.930817	0.383836	3.069894	8.740447	0.158053	0.368331
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	2017.000000	1.000000	1.000000	0.000000	0.000000
25%	2.000000	0.000000	0.000000	1.000000	0.000000	17.000000	2018.000000	5.000000	8.000000	0.000000	0.000000
50%	2.000000	0.000000	1.000000	2.000000	0.000000	57.000000	2018.000000	8.000000	16.000000	0.000000	0.000000
75%	2.000000	0.000000	2.000000	3.000000	0.000000	126.000000	2018.000000	10.000000	23.000000	0.000000	0.000000
max	4.000000	10.000000	7.000000	17.000000	1.000000	443.000000	2018.000000	12.000000	31.000000	1.000000	13.000000

A hotel will not rent out their rooms for free, unless there is a special offer in place. However, there is no indication that the hotels are giving special offers, thus making this number impossible. Therefore, 0 in average price per room will be treated as missing value.

no_of_previous_bookings_not_canceled	avg_price_per_room	no_of_special_requests
36275.000000	36275.000000	36275.000000
0.153411	103.423539	0.619655
1.754171	35.089424	0.786236
0.000000	0.000000	0.000000
0.000000	80.300000	0.000000
0.000000	99.450000	0.000000
0.000000	120.000000	1.000000
58.000000	540.000000	5.000000

PROFILING DESCRIPTIVE STATISTICS

```
# Finding and treating missing values  
data.loc[data['avg_price_per_room']==0]
```



545 rows × 19 columns

Treatment:

```
data['avg_price_per_room'] = data['avg_price_per_room'].replace([0],data['avg_price_per_room'].median())
```



```
data['avg_price_per_room'].describe()
```

```
count    36275.000000  
mean      104.917688  
std       32.688889  
min       0.500000  
25%       81.000000  
50%       99.450000  
75%      120.000000  
max       540.000000
```

```
Name: avg_price_per_room, dtype: float64
```


PROFILING DESCRIPTIVE STATISTICS

```
#Object data profiling  
data.describe(include='O')
```

	Booking_ID	type_of_meal_plan	room_type_reserved	market_segment_type	booking_status
count	36275	36275	36275	36275	36275
unique	36275	4	7	5	2
top	INN00001	Meal Plan 1	Room_Type 1	Online	Not_Canceled
freq	1	27835	28130	23214	24390

CHECKING DUPLICATED DATA

```
data[data.duplicated()]
```

Booking_ID	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_nights	type_of_meal_plan	required_car_parking_space	room_type_reserved
------------	--------------	----------------	----------------------	-------------------	-------------------	----------------------------	--------------------

There is no duplicated entry.

CHECKING THE DATA BALANCE

```
data.groupby('booking_status').size()
```

```
booking_status  
Canceled      11885  
Not_Canceled  24390  
dtype: int64
```

```
data['booking_status'].value_counts(normalize=True)
```

```
Not_Canceled    0.672364  
Canceled        0.327636  
Name: booking_status, dtype: float64
```

The data is not balance. 67% did not cancel, and 33% canceled.

SIMPLIFYING FEATURES NAME

```
data.rename(columns={'no_of_adults':'adults'},inplace=True)
data.rename(columns={'no_of_children':'children'},inplace=True)
data.rename(columns={'no_of_weekend_nights':'weekend_nights'},inplace=True)
data.rename(columns={'no_of_week_nights':'weekday_nights'},inplace=True)
data.rename(columns={'type_of_meal_plan':'meal_plan'},inplace=True)
data.rename(columns={'no_of_weekend_nights':'weekend_nights'},inplace=True)
data.rename(columns={'no_of_special_requests':'special_requests'},inplace=True)
data.rename(columns={'market_segment_type':'market_type'},inplace=True)
data.rename(columns={'special_requests':'request_num'},inplace=True)

data.rename(columns={'required_car_parking_space':'parking_space'},inplace=True)

data.rename(columns={'room_type_reserved':'room_type'},inplace=True)
data.rename(columns={'avg_price_per_room':'avg_room_price'},inplace=True)
```

SIMPLIFYING FEATURES NAME


#	Column
0	Booking_ID
1	no_of_adults
2	no_of_children
3	no_of_weekend_nights
4	no_of_week_nights
5	type_of_meal_plan
6	required_car_parking_space
7	room_type_reserved
8	lead_time
9	arrival_year
10	arrival_month
11	arrival_date
12	market_segment_type
13	repeated_guest
14	no_of_previous_cancellations
15	no_of_previous_bookings_not_canceled
16	avg_price_per_room
17	no_of_special_requests
18	booking_status

Before vs After

#	Column
0	Booking_ID
1	adults
2	children
3	weekend_nights
4	weekday_nights
5	meal_plan
6	parking_space
7	room_type
8	lead_time
9	arrival_year
10	arrival_month
11	arrival_date
12	market_type
13	repeated_guest
14	no_of_previous_cancellations
15	no_of_previous_bookings_not_canceled
16	avg_room_price
17	request_num
18	booking_status

CHANGING THE TARGET INTO NUMERIC

```
data['booking_status'] = data['booking_status'].replace(['Not_Canceled', 'Canceled'],[0,1])
```

booking_status		booking_status
Not_Canceled		0
Not_Canceled		0
Canceled		1
Canceled		1
Canceled		1

FEATURE SCALLING AND ENCODING

```
#Importing Scaler and Encoder  
from sklearn.preprocessing import RobustScaler  
from sklearn.preprocessing import OneHotEncoder
```

Robust Scaler is used as we leave the outliers in the data. **One Hot Encoder** itself is used because all the categorical feature only serve as category without hierarchy.

```
rscaler = RobustScaler()  
  
ohc = OneHotEncoder(handle_unknown = 'ignore')
```

FEATURE SCALLING AND ENCODING

```
data.drop(columns=['Booking_ID'], axis=1, inplace=True)
```

Booking ID is dropped as it will not be used in Machine Learning.

```
# Copying the dataset for scalling  
data_scaled = data.copy()
```

```
data_scaled.head()
```

	adults	children	weekend_nights	weekday_nights	meal_plan	parking_space	room_type	lead_time
0	2	0	1	2	Meal Plan 1	0	Room_Type 1	224
1	2	0	2	3	Not Selected	0	Room_Type 1	5
2	1	0	2	1	Meal Plan 1	0	Room_Type 1	1
3	2	0	0	2	Meal Plan 1	0	Room_Type 1	211
4	2	0	1	1	Not Selected	0	Room_Type 1	48

FEATURE SCALING

```
# Scaling the numerical features
data_scaled['adults'] = rscaler.fit_transform(data_scaled[['adults']])
data_scaled['children'] = rscaler.fit_transform(data_scaled[['children']])
data_scaled['weekend_nights'] = rscaler.fit_transform(data_scaled[['weekend_nights']])
data_scaled['weekday_nights'] = rscaler.fit_transform(data_scaled[['weekday_nights']])
data_scaled['lead_time'] = rscaler.fit_transform(data_scaled[['lead_time']])
data_scaled['no_of_previous_cancellations'] = rscaler.fit_transform(data_scaled[['no_of_previous_cancellations']])
data_scaled['no_of_previous_bookings_not_canceled'] = rscaler.fit_transform(data_scaled[['no_of_previous_bookings_not_canceled']])
data_scaled['avg_room_price'] = rscaler.fit_transform(data_scaled[['avg_room_price']])
data_scaled['request_num'] = rscaler.fit_transform(data_scaled[['request_num']])
```

data_scaled.head()

	adults	children	weekend_nights	weekday_nights	meal_plan	parking_space	room_type	lead_time
0	0.0	0.0	0.0	0.0	Meal Plan 1	0	Room_Type 1	1.532110
1	0.0	0.0	0.5	0.5	Not Selected	0	Room_Type 1	-0.477064
2	-1.0	0.0	0.5	-0.5	Meal Plan 1	0	Room_Type 1	-0.513761
3	0.0	0.0	-0.5	0.0	Meal Plan 1	0	Room_Type 1	1.412844
4	0.0	0.0	0.0	-0.5	Not Selected	0	Room_Type 1	-0.082569

FEATURE ENCODING

```
col = sorted(data['meal_plan'].unique().tolist()) + sorted(data['room_type'].unique().tolist()) + sorted(data['market_type'].unique().tolist())
```

```
#Encoding
```

```
enc_df = pd.DataFrame(ohc.fit_transform(data_scaled[['meal_plan', 'room_type', 'market_type']]).toarray(), columns=col)
```

```
enc_df.head()
```

	Meal Plan 1	Meal Plan 2	Meal Plan 3	Not Selected	Room_Type 1	Room_Type 2	Room_Type 3	Room_Type 4	Room_Type 5	Room_Type 6	Room_Type 7	Aviation	Complementary	Corporate	Offline	Online
0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
2	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
3	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0

FEATURE SCALLING AND ENCODING

```
#Joining the scaled dataframe with encoded dataframe
```

```
data_scaled = data_scaled.join(enc_df)
```

```
data_scaled.head()
```

	adults	children	weekend_nights	weekday_nights	meal_plan	parking_space	room_type	lead_time	arrival_year	arrival_month	...	Room_Type_3	Room_Type_4	Room_Type_5	Room_Type_6	Room_Type_7	Aviation
0	0.0	0.0	0.0	0.0	Meal Plan 1	0	Room_Type 1	1.532110	2017	10	...	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.5	0.5	Not Selected	0	Room_Type 1	-0.477064	2018	11	...	0.0	0.0	0.0	0.0	0.0	0.0
2	-1.0	0.0	0.5	-0.5	Meal Plan 1	0	Room_Type 1	-0.513761	2018	2	...	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	-0.5	0.0	Meal Plan 1	0	Room_Type 1	1.412844	2018	5	...	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	-0.5	Not Selected	0	Room_Type 1	-0.082569	2018	4	...	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 34 columns

FEATURE SCALLING AND ENCODING

```
In [298]: # Dropping the unused features
data_scaled.drop(columns=['meal_plan','room_type','market_type', 'arrival_year','arrival_month','arrival_date'],axis=1,inplace=True)
```

```
In [301]: # joining dummies to the dataframe
final_data=data_scaled.copy()
```

```
0  adults      36275 non-null float64
1  children    36275 non-null float64
2  weekend_nights  36275 non-null float64
3  weekday_nights  36275 non-null float64
4  parking_space  36275 non-null int64
5  lead_time    36275 non-null float64
6  repeated_guest  36275 non-null int64
7  no_of_previous_cancellations  36275 non-null float64
8  no_of_previous_bookings_not_canceled  36275 non-null float64
9  avg_room_price  36275 non-null float64
10 request_num  36275 non-null float64
11 booking_status  36275 non-null int64
12 Meal Plan 1    36275 non-null float64
13 Meal Plan 2    36275 non-null float64
14 Meal Plan 3    36275 non-null float64
15 Not Selected   36275 non-null float64
16 Room_Type 1    36275 non-null float64
17 Room_Type 2    36275 non-null float64
18 Room_Type 3    36275 non-null float64
19 Room_Type 4    36275 non-null float64
20 Room_Type 5    36275 non-null float64
21 Room_Type 6    36275 non-null float64
22 Room_Type 7    36275 non-null float64
23 Aviation       36275 non-null float64
24 Complementary  36275 non-null float64
25 Corporate      36275 non-null float64
26 Offline        36275 non-null float64
27 Online         36275 non-null float64
```

All the features is either integer or float.

TRAIN TEST SPLIT

```
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from collections import Counter
```

#Splitting X and Y

```
X = final_data.drop(columns=['booking_status'],axis=1)
y = final_data['booking_status']
```

#split train-rest data

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25,random_state=0)
```

X_train.shape

(27206, 28)

y_train.shape

(27206,)

BALANCING DATA

```
print('Before Oversampling: ',Counter(y_train))  
#defining smote  
SMOTE = SMOTE(random_state=0)  
  
#fit and apply the transform  
X_train_smote, y_train_smote = SMOTE.fit_resample(X_train,y_train)  
  
#summarize class distribution  
print('After Oversampling: ',Counter(y_train_smote))
```

Before Oversampling: Counter({0: 18328, 1: 8878})

After Oversampling: Counter({0: 18328, 1: 18328})

MACHINE LEARNING MODELING

(Decision Tree Classifier)

```
# Decision Tree without resampling  
dt = DecisionTreeClassifier(random_state=0)  
dt.fit(X_train, y_train)
```

```
DecisionTreeClassifier  
DecisionTreeClassifier(random_  
state=0)
```

```
# Decision Tree with resampling  
dt2 = DecisionTreeClassifier(random_state=0)  
dt2.fit(X_train_smote, y_train_smote)
```

```
DecisionTreeClassifier  
DecisionTreeClassifier(random_  
state=0)
```



```
y_pred_dt = dt.predict(X_test)  
print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
0	0.89	0.90	0.89	6062
1	0.79	0.77	0.78	3007
accuracy			0.85	9069
macro avg	0.84	0.83	0.83	9069
weighted avg	0.85	0.85	0.85	9069



```
y_pred_dt2 = dt2.predict(X_test)  
print(classification_report(y_test, y_pred_dt2))
```

	precision	recall	f1-score	support
0	0.89	0.87	0.88	6062
1	0.75	0.78	0.77	3007
accuracy			0.84	9069
macro avg	0.82	0.83	0.83	9069
weighted avg	0.85	0.84	0.84	9069

MACHINE LEARNING MODELING

(Random Forest Classifier)

```
# Random Forest without resampling  
rf1 = RandomForestClassifier(random_state=0)  
rf1.fit(X_train,y_train)
```

RandomForestClassifier

```
RandomForestClassifier(random_  
state=0)
```

```
# Random Forest with resampling  
rf2 = RandomForestClassifier(random_state=0)  
rf2.fit(X_train_smote,y_train_smote)
```

RandomForestClassifier

```
RandomForestClassifier(random_  
state=0)
```

```
y_pred_rf1 = rf1.predict(X_test)  
print(classification_report(y_test, y_pred_rf1))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.92	6062
1	0.86	0.78	0.82	3007
accuracy			0.88	9069
macro avg	0.88	0.86	0.87	9069
weighted avg	0.88	0.88	0.88	9069

```
y_pred_rf2 = rf2.predict(X_test)  
print(classification_report(y_test, y_pred_rf2))
```

	precision	recall	f1-score	support
0	0.90	0.91	0.91	6062
1	0.82	0.80	0.81	3007
accuracy			0.88	9069
macro avg	0.86	0.86	0.86	9069
weighted avg	0.88	0.88	0.88	9069

MACHINE LEARNING MODELING (KNN)

```
# KNN without resampling  
knn = KNeighborsClassifier()  
knn.fit(X_train, y_train)
```

```
▼ KNeighborsClassifier  
KNeighborsClassifi  
er()
```

```
# KNN with resampling  
knn2 = KNeighborsClassifier()  
knn2.fit(X_train_smote, y_train_smote)
```

```
▼ KNeighborsClassifier  
KNeighborsClassifi  
er()
```




```
y_pred_knn = knn.predict(X_test)  
print(classification_report(y_test, y_pred_knn))
```

	precision	recall	f1-score	support
0	0.88	0.91	0.89	6062
1	0.80	0.74	0.77	3007
accuracy			0.85	9069
macro avg	0.84	0.82	0.83	9069
weighted avg	0.85	0.85	0.85	9069

```
y_pred_knn2 = knn2.predict(X_test)  
print(classification_report(y_test, y_pred_knn2))
```

	precision	recall	f1-score	support
0	0.90	0.85	0.87	6062
1	0.72	0.81	0.76	3007
accuracy			0.83	9069
macro avg	0.81	0.83	0.82	9069
weighted avg	0.84	0.83	0.84	9069



MACHINE LEARNING MODELING

(Support Vector Classifier, Linear Kernel)

SVCL without resampling

```
svc = SVC(kernel='linear', random_state=0)
svc.fit(X_train, y_train)
```

SVC

```
SVC(kernel='linear', random_s  
tate=0)
```

SVCL with resampling

```
svc2 = SVC(kernel='linear', random_state=0)
svc2.fit(X_train_smote, y_train_smote)
```


SVC

```
SVC(kernel='linear', random_s  
tate=0)
```



```
y_pred_svc = svc.predict(X_test)
print(classification_report(y_test, y_pred_svc))
```

	precision	recall	f1-score	support
0	0.82	0.90	0.86	6062
1	0.75	0.59	0.66	3007
accuracy			0.80	9069
macro avg	0.78	0.75	0.76	9069
weighted avg	0.79	0.80	0.79	9069



```
y_pred_svc2 = svc2.predict(X_test)
print(classification_report(y_test, y_pred_svc2))
```

	precision	recall	f1-score	support
0	0.87	0.77	0.82	6062
1	0.63	0.77	0.69	3007
accuracy			0.77	9069
macro avg	0.75	0.77	0.76	9069
weighted avg	0.79	0.77	0.78	9069

MACHINE LEARNING MODELING

(Support Vector Classifier, RBF Kernel)

```
# SVCR without resampling  
svc3 = SVC(kernel='rbf', random_state=0)  
svc3.fit(X_train, y_train)
```

SVC

```
SVC(random_state=0)
```

```
# SVCR with resampling  
svc4 = SVC(kernel='rbf', random_state=0)  
svc4.fit(X_train_smote, y_train_smote)
```


SVC

```
SVC(random_state=0)
```



```
y_pred_svc3 = svc3.predict(X_test)  
print(classification_report(y_test, y_pred_svc3))
```

	precision	recall	f1-score	support
0	0.83	0.94	0.88	6062
1	0.82	0.61	0.70	3007
accuracy			0.83	9069
macro avg	0.83	0.77	0.79	9069
weighted avg	0.83	0.83	0.82	9069



```
y_pred_svc4 = svc4.predict(X_test)  
print(classification_report(y_test, y_pred_svc4))
```

	precision	recall	f1-score	support
0	0.89	0.83	0.86	6062
1	0.69	0.80	0.74	3007
accuracy			0.82	9069
macro avg	0.79	0.81	0.80	9069
weighted avg	0.83	0.82	0.82	9069

MACHINE LEARNING MODELING

(Best Performing Model)


```
y_pred_rf1 = rf1.predict(X_test)
print(classification_report(y_test, y_pred_rf1))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.92	6062
1	0.86	0.78	0.82	3007
accuracy			0.88	9069
macro avg	0.88	0.86	0.87	9069
weighted avg	0.88	0.88	0.88	9069

Random Forest Classifier without resampling

CHECKING FEATURE IMPORTANCE


In the model earlier, it still uses a lot of features. I aim to check the importance of each features so I can reduce it thus making the model work more efficiently.



CHECKING FEATURE IMPORTANCE

```
# Checking Feature Importance
features = X.columns
importances = rf1.feature_importances_
indices = np.argsort(importances)[-15:]
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='b', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

CHECKING FEATURE IMPORTANCE



RETRAINING AFTER CHECKING FEATURE IMPORTANCE

```
#Splitting X and Y
```

```
# X = final_data['adults','children','lead_time','avg_room_price','request_num','weekday_nights','weekend_nights','Online','Offline']
```

```
X = final_data.drop(columns=['parking_space','repeated_guest','no_of_previous_cancellations','no_of_previous_bookings_not_canceled',  
                             'Meal Plan 1','Meal Plan 2','Meal Plan 3','Not Selected','Room_Type 1','Room_Type 2','Room_Type 3','Room_Type 4',  
                             'Room_Type 5','Room_Type 6','Room_Type 7','Aviation','Complementary','Corporate','booking_status'], axis=1)
```

```
y = final_data['booking_status']
```

```
#split train-test data
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25,random_state=0)
```

RETRAINING AFTER CHECKING FEATURE IMPORTANCE

```
y_pred_rf1 = rf1.predict(X_test)
print(classification_report(y_test, y_pred_rf1))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.92	6062
1	0.86	0.78	0.82	3007
accuracy			0.88	9069
macro avg	0.88	0.86	0.87	9069
weighted avg	0.88	0.88	0.88	9069

BEFORE

```
y_pred_rf3 = rf3.predict(X_test)
print(classification_report(y_test, y_pred_rf3))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	6062
1	0.85	0.76	0.80	3007
accuracy			0.88	9069
macro avg	0.87	0.85	0.86	9069
weighted avg	0.87	0.88	0.87	9069

AFTER

There is not much difference after we reduce the number of features used. This performance can be improved with hyperparameter tuning. In this project, I am using the Grid Search CV.

HYPERPARAMETER TUNING

```
from sklearn.model_selection import GridSearchCV
```

```
# Define the parameter grid to search over
```

```
param_grid = {  
    'n_estimators': [10, 50, 100, 200],  
    'max_depth': [None, 5, 10, 20],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4],  
}
```

```
# Create the GridSearchCV object
```

```
grid_search = GridSearchCV(rf3, param_grid=param_grid, cv=5)
```

```
# Fit the GridSearchCV object to the data
```

```
grid_search.fit(X, y)
```

```
# Print the best hyperparameters and the corresponding score
```

```
print("Best hyperparameters:", grid_search.best_params_)
```

```
print("Best score:", grid_search.best_score_)
```

```
Best hyperparameters: {'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200}  
Best score: 0.8831426602343211
```

HYPERPARAMETER TUNING

```
grid_search.best_params_
```

```
{'max_depth': 20,  
 'min_samples_leaf': 1,  
 'min_samples_split': 2,  
 'n_estimators': 200}
```

```
grid_search.best_estimator_
```

▼

RandomForestClassifier

```
RandomForestClassifier(max_depth=20, n_estimators=200, r  
andom_state=0)
```

HYPERPARAMETER TUNING RESULT

```
y_pred_rf3 = rf3.predict(X_test)
print(classification_report(y_test, y_pred_rf3))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	6062
1	0.85	0.76	0.80	3007
accuracy			0.88	9069
macro avg	0.87	0.85	0.86	9069
weighted avg	0.87	0.88	0.87	9069

BEFORE

```
best_grid = grid_search.best_estimator_

y_pred_grid = best_grid.predict(X_test)
print(classification_report(y_test, y_pred_grid))
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	6062
1	0.97	0.92	0.95	3007
accuracy			0.97	9069
macro avg	0.97	0.95	0.96	9069
weighted avg	0.97	0.97	0.97	9069

AFTER

After tuning, we can see that the performance of the model is significantly improved.

THANK YOU